

Data Wrangling with dplyr Part 1

Week 3 Friday

Miles Chen

Department of Statistics and Data Science



Section 1

dplyr

dplyr

dplyr is a powerful package for data manipulation and a core component of the tidyverse.

You can load dplyr specifically with:

```
library(dplyr)
```

Or load the entire tidyverse, which includes dplyr and other related packages:

```
library(tidyverse)
```

dplyr

dplyr is a package in the tidyverse that provides a consistent set of functions that help you solve data manipulation challenges.

cheat sheet

<https://github.com/rstudio/cheatsheets/blob/main/data-transformation.pdf>

The starwars Dataset

The `starwars` dataset is included with `dplyr`. It contains information about various Star Wars characters from the first seven Star Wars movies. This dataset is structured as a tibble (a kind of data frame).

`starwars`

```
# A tibble: 87 x 14
  name      height  mass hair_color skin_color eye_color birth_year sex   gender homeworld species
  <chr>     <int> <dbl> <chr>       <chr>       <chr>       <dbl> <chr> <chr>       <chr>       <chr>
1 Luke Skywalker 172    77 blond      fair        blue         19   male   masculin Tatooine Human
2 C-3PO          167    75 <NA>       gold        yellow       112  none   masculin Tatooine Droid 
3 R2-D2          96     32 <NA>       white, bl~ red         33   none   masculin Naboo   Droid 
4 Darth Vader   202    136 none       white        yellow       41.9 male   masculin Tatooine Human
5 Leia Organa    150    49 brown      light       brown        19   female feminin Alderaan Human
6 Owen Lars     178    120 brown, gr~ light       blue         52   male   masculin Tatooine Human
7 Beru Whitesun 165    75 brown      light       blue         47   female feminin Tatooine Human
8 R5-D4          97     32 <NA>       white, red red        NA   none   masculin Tatooine Droid 
9 Biggs Darklighter 183   84 black      light       brown        24   male   masculin Tatooine Human
10 Obi-Wan Kenobi 182   77 auburn, w~ fair        blue-gray      57   male   masculin Stewjon Human
# i 77 more rows
# i 3 more variables: films <list>, vehicles <list>, starships <list>
```

glimpse()

The glimpse() function (part of dplyr) provides a structured overview of a dataset. It displays column names, data types, and a preview of values. Similar to str().

```
glimpse(starwars)
```

```
Rows: 87
Columns: 14
$ name      <chr> "Luke Skywalker", "C-3PO", "R2-D2", "Darth Vader", "Leia Organa", "Owen Lars", ~
$ height     <int> 172, 167, 96, 202, 150, 178, 165, 97, 183, 182, 188, 180, 228, 180, 173, 175, 1~
$ mass       <dbl> 77.0, 75.0, 32.0, 136.0, 49.0, 120.0, 75.0, 32.0, 84.0, 77.0, 84.0, NA, 112.0, ~
$ hair_color <chr> "blond", NA, NA, "none", "brown", "brown", "grey", "brown", NA, "black", "auburn", ~
$ skin_color <chr> "fair", "gold", "white", "blue", "white", "light", "light", "light", "white", "red", ~
$ eye_color  <chr> "blue", "yellow", "red", "yellow", "brown", "blue", "blue", "red", "brown", "bl~
$ birth_year <dbl> 19.0, 112.0, 33.0, 41.9, 19.0, 52.0, 47.0, NA, 24.0, 57.0, 41.9, 64.0, 200.0, 2~
$ sex        <chr> "male", "none", "none", "male", "female", "male", "female", "none", "male", "ma~
$ gender     <chr> "masculine", "masculine", "masculine", "masculine", "feminine", "masculine", "f~
$ homeworld  <chr> "Tatooine", "Tatooine", "Naboo", "Tatooine", "Alderaan", "Tatooine", "Tatooine", ~
$ species    <chr> "Human", "Droid", "Droid", "Human", "Human", "Human", "Human", "Droid", "Human", ~
$ films      <list> <"A New Hope", "The Empire Strikes Back", "Return of the Jedi", "Revenge of th~
$ vehicles   <list> <"Snowspeeder", "Imperial Speeder Bike">, <>, <>, <>, "Imperial Speeder Bike", ~
$ starships   <list> <"X-wing", "Imperial shuttle">, <>, <>, "TIE Advanced x1", <>, <>, <>, <>, "X-~
```

The nycflights23 Dataset

Contains information from flights in New York City in the year 2023.

```
library(nycflights23)  
data(flights)
```

```
glimpse(flights)
```

Rows: 435,352

Columns: 19

flights

```
# A tibble: 435,352 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time arr_delay carrier
  <int> <int> <int>     <int>           <int>     <dbl>     <int>           <int>     <dbl> <chr>
1 2023     1     1         1            2038      203       328            3        205  UA
2 2023     1     1        18            2300       78       228          135        53  DL
3 2023     1     1        31            2344       47       500          426        34  B6
4 2023     1     1        33            2140      173       238          2352       166  B6
5 2023     1     1        36            2048      228       223          2252       211  UA
6 2023     1     1       503            500        3       808          815        -7  AA
7 2023     1     1       520            510       10       948          949        -1  B6
8 2023     1     1       524            530       -6       645          710       -25  AA
9 2023     1     1       537            520       17       926          818        68  UA
10 2023    1     1       547            545        2       845          852        -7  NK
# i 435,342 more rows
# i 9 more variables: flight <int>, tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>,
#   distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dttm>
```

The Pipe

R has a pipe operator `|>` (older versions may appear as `%>%`)

The pipe `|>` takes the result of what is in front of the pipe and inserts it as the first argument in the function that comes after the pipe. `x |> f(y)` is equivalent to `f(x, y)`.

The easiest way to pronounce the pipe is “then”.

The Pipe

Instead of writing:

```
temp_data1 <- func1(original_data, options1)
temp_data2 <- func2(temp_data1, options2)
final_data <- func3(temp_data2, options3)
```

We can string them together with pipes:

```
final_data <- original_data |>
  func1(options1) |>
  func2(options2) |>
  func3(options3)
```

Shortcut to insert the pipe

Use this keyboard shortcut to quickly insert the pipe operator in RStudio:

Windows/Linux: CTRL + SHIFT + M

Mac: CMD + SHIFT + M

Depending on your settings, RStudio may insert the pipe function (%>%) from the package `magrittr`, which works very similarly.

Section 2

dplyr basics

dplyr basics

The functions in dplyr each perform a different task and all have the following in common:

- The first argument is always a data frame.
- The subsequent arguments usually describe which columns to operate on.
- The output is always a new data frame.

Each function does one thing, so many data wrangling problems will require combining multiple functions. We accomplish with the pipe.

Selecting Columns with `select()`

The `select()` function lets you choose columns from a dataset. When using `select()`, you do not need to use quotes around column names unless they contain spaces or special characters.

```
select(starwars, name, homeworld, species, films)
```

```
# A tibble: 87 x 4
  name          homeworld species   films
  <chr>        <chr>     <chr>    <list>
1 Luke Skywalker Tatooine  Human    <chr [5]>
2 C-3PO          Tatooine  Droid    <chr [6]>
3 R2-D2          Naboo     Droid    <chr [7]>
4 Darth Vader   Tatooine  Human    <chr [4]>
5 Leia Organa   Alderaan  Human    <chr [5]>
6 Owen Lars     Tatooine  Human    <chr [3]>
7 Beru Whitesun Lars Tatooine  Human    <chr [3]>
8 R5-D4          Tatooine  Droid    <chr [1]>
9 Biggs Darklighter Tatooine  Human    <chr [1]>
10 Obi-Wan Kenobi Stewjon   Human    <chr [6]>
# i 77 more rows
```

Selecting Columns with `select()`

- Use a negative sign (-) to **deselect** columns.

```
starwars |>  
  select(-eye_color, -birth_year) |>  
  print(n = 6)
```

```
# A tibble: 87 x 12  
   name    height  mass hair_color skin_color sex  gender homeworld species films vehicles starships  
   <chr>     <int> <dbl> <chr>      <chr> <chr> <chr> <chr> <list> <list> <list>  
 1 Luke S~    172     77 blond     fair   male  masculin Tatooine Human  <chr> <chr> <chr> [2]>  
 2 C-3PO       167     75 <NA>      gold   none  masculin Tatooine Droid  <chr> <chr> <chr> [0]>  
 3 R2-D2       96      32 <NA>      white, bl~ none  masculin Naboo   Droid  <chr> <chr> <chr> [0]>  
 4 Darth ~     202     136 none      white   male  masculin Tatooine Human  <chr> <chr> <chr> [1]>  
 5 Leia O~     150     49 brown     light   femal feminin Alderaan Human  <chr> <chr> <chr> [0]>  
 6 Owen L~     178     120 brown, gr~ light   male  masculin Tatooine Human  <chr> <chr> <chr> [0]>  
# i 81 more rows
```

This removes the `name`, `eye_color`, and `birth_year` columns while keeping all others.

Selecting a Range of Columns

- Use **colon notation (:) to select a continuous range of columns.**

```
starwars |>  
  select(name:eye_color) |>  
  print(n = 6)
```

```
# A tibble: 87 x 6  
  name      height  mass hair_color skin_color eye_color  
  <chr>     <int> <dbl> <chr>       <chr>      <chr>  
1 Luke Skywalker    172    77 blond      fair       blue  
2 C-3PO            167    75 <NA>       gold       yellow  
3 R2-D2             96     32 <NA>      white, blue red  
4 Darth Vader      202   136 none       white       yellow  
5 Leia Organa       150     49 brown      light      brown  
6 Owen Lars         178   120 brown, grey light      blue  
# i 81 more rows
```

This keeps only the columns from `name` through `eye_color`, in the order they appear in the dataset.

Special Selection Functions in dplyr

dplyr has special selection functions. See `?tidyselect::select_helpers`

- `starts_with("x")`, `ends_with("x")`, `contains("x")` - picks columns that match a string
- `matches(regex)` picks columns that match a regular expression
- `any_of(c("colA", "colB"))` ignores missing names
- `everything()` Select all (remaining) columns
- `num_range("x", 1:5)` Select columns named something like `x1`, `x2`, `x3`, `x4`, `x5`
- `one_of(name_vector)` Select columns where the names are stored in a vector

Example of Special Selection Functions

Select name and all columns that end with "color":

```
flights |>  
  select(contains("time"))
```

```
# A tibble: 435,352 x 6  
  dep_time sched_dep_time arr_time sched_arr_time air_time time_hour  
    <int>        <int>     <int>        <int>      <dbl> <dttm>  
1       1          2038      328           3       367 2023-01-01 20:00:00  
2      18          2300      228          135      108 2023-01-01 23:00:00  
3      31          2344      500          426      190 2023-01-01 23:00:00  
4      33          2140      238          2352     108 2023-01-01 21:00:00  
5      36          2048      223          2252      80  2023-01-01 20:00:00  
6     503           500      808          815      154 2023-01-01 05:00:00  
7     520           510      948          949      192 2023-01-01 05:00:00  
8     524           530      645          710      119 2023-01-01 05:00:00  
9     537           520      926          818      258 2023-01-01 05:00:00  
10    547           545      845          852      157 2023-01-01 05:00:00  
# i 435,342 more rows
```

Examples of Special Selection Functions

```
flights |>
  select(starts_with("dep"), starts_with("arr"), everything())

# A tibble: 435,352 x 19
  dep_time dep_delay arr_time arr_delay year month   day sched_dep_time sched_arr_time carrier
    <int>     <dbl>    <int>     <dbl> <int> <int> <int>          <int>          <int> <chr>
1       1      203     328      205  2023     1     1        2038         3  UA
2      18       78     228       53  2023     1     1        2300        135  DL
3      31       47     500       34  2023     1     1        2344        426  B6
4      33      173     238      166  2023     1     1        2140       2352  B6
5      36      228     223      211  2023     1     1        2048       2252  UA
6     503        3     808      -7  2023     1     1         500        815  AA
7     520       10     948      -1  2023     1     1         510        949  B6
8     524       -6     645     -25  2023     1     1         530        710  AA
9     537       17     926      68  2023     1     1         520        818  UA
10    547        2     845      -7  2023     1     1         545        852  NK
# i 435,342 more rows
# i 9 more variables: flight <int>, tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>,
#   distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dttm>
```

Rename while selecting

Rename while selecting (new = old):

```
flights |>  
  select(dep = dep_time, arr = arr_time, everything())
```

```
# A tibble: 435,352 x 19  
  dep    arr   year month   day sched_dep_time dep_delay sched_arr_time arr_delay carrier flight  
  <int> <int> <int> <int> <int>       <int>     <dbl>       <int>     <dbl> <chr>    <int>  
1     1    328  2023     1     1        2038      203          3      205  UA      628  
2    18    228  2023     1     1        2300       78         135      53  DL      393  
3    31    500  2023     1     1        2344       47         426      34  B6      371  
4    33    238  2023     1     1        2140      173        2352      166 B6     1053  
5    36    223  2023     1     1        2048      228        2252      211 UA      219  
6   503    808  2023     1     1        500        3        815      -7  AA      499  
7   520    948  2023     1     1        510       10        949      -1  B6      996  
8   524    645  2023     1     1        530       -6        710     -25 AA     981  
9   537    926  2023     1     1        520       17        818      68  UA      206  
10  547    845  2023     1     1        545        2        852      -7  NK      225  
# i 435,342 more rows  
# i 8 more variables: tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,  
#   hour <dbl>, minute <dbl>, time_hour <dttm>
```

Selecting Columns with a Variable

You can dynamically select columns stored in a character vector using `all_of()` or `any_of()`.

- `all_of(vars)` – Selects columns but **errors** out if any are missing.
- `any_of(vars)` – Selects columns without error if some are missing.

```
vars <- c("month", "day", "carrier", "flight", "origin", "dest")
flights |>
  select(all_of(vars))
```

```
# A tibble: 435,352 x 6
  month   day carrier flight origin dest
  <int> <int> <chr>    <int> <chr>  <chr>
1     1     1 UA        628 EWR    SMF
2     1     1 DL        393 JFK    ATL
3     1     1 B6        371 JFK    BQN
4     1     1 B6       1053 JFK    CHS
5     1     1 UA        219 EWR    DTW
6     1     1 AA        499 EWR    MIA
7     1     1 B6        996 JFK    BQN
8     1     1 AA        981 EWR    ORD
9     1     1 UA        206 EWR    IAH
10    1     1 NK        225 EWR    FLL
```

Selecting a column and using the values with pull()

If you want a vector with the values, use pull()

```
starwars |>  
  select(name) |>  
  pull()
```

[1] "Luke Skywalker"	"C-3PO"	"R2-D2"	"Darth Vader"
[5] "Leia Organa"	"Owen Lars"	"Beru Whitesun Lars"	"R5-D4"
[9] "Biggs Darklighter"	"Obi-Wan Kenobi"	"Anakin Skywalker"	"Wilhuff Tarkin"
[13] "Chewbacca"	"Han Solo"	"Greedo"	"Jabba Desilijic Tiure"
[17] "Wedge Antilles"	"Jek Tono Porkins"	"Yoda"	"Palpatine"
[21] "Boba Fett"	"IG-88"	"Bossk"	"Lando Calrissian"
[25] "Lobot"	"Ackbar"	"Mon Mothma"	"Arvel Crynyd"
[29] "Wicket Systri Warrick"	"Nien Nunb"	"Qui-Gon Jinn"	"Nute Gunray"
[33] "Finis Valorum"	"Padmé Amidala"	"Jar Jar Binks"	"Roos Tarpals"
[37] "Rugor Nass"	"Ric Olié"	"Watto"	"Sebulba"
[41] "Quarsh Panaka"	"Shmi Skywalker"	"Darth Maul"	"Bib Fortuna"
[45] "Ayla Secura"	"Ratts Tyerel"	"Dud Bolt"	"Gasgano"
[49] "Ben Quadinaros"	"Mace Windu"	"Ki-Adi-Mundi"	"Kit Fisto"
[53] "Eeth Koth"	"Adi Gallia"	"Saesee Tiin"	"Yarael Poof"
[57] "Plo Koon"	"Mas Amedda"	"Gregar Typho"	"Cordé"
[61] "Cliegg Lars"	"Poggle the Lesser"	"Luminara Unduli"	"Barriss Offee"
[65] "Dormé"	"Dooku"	"Bail Prestor Organa"	"Jango Fett"
[69] "Zannah"	"Dexter Jettster"	"Lama Su"	"Taun We"

Filtering Rows with filter()

The `filter()` function allows you to **select rows** based on conditions. It keeps only the rows where the condition evaluates to TRUE.

Conditions must return a **logical vector** (TRUE/FALSE) with a length equal to the number of rows in the dataset.

```
starwars |>  
  filter(species == "Droid")
```

```
# A tibble: 6 x 14  
  name    height  mass hair_color skin_color eye_color birth_year sex   gender homeworld species films  
  <chr>   <int> <dbl> <chr>     <chr>      <chr>       <dbl> <chr> <chr>    <chr>    <chr>  <list>  
1 C-3PO    167    75 <NA>      gold       yellow      112  none  mascul~ Tatooine  Droid  <chr>  
2 R2-D2     96     32 <NA>      white, bl~ red        33  none  mascul~ Naboo   Droid  <chr>  
3 R5-D4     97     32 <NA>      white, red red      NA  none  mascul~ Tatooine Droid  <chr>  
4 IG-88    200    140 none      metal       red        15  none  mascul~ <NA>   Droid  <chr>  
5 R4-P~     96     NA none      silver, r~ red, blue  NA  none  femin~ <NA>   Droid  <chr>  
6 BB8      NA     NA none      none        black      NA  none  mascul~ <NA>   Droid  <chr>  
# i 2 more variables: vehicles <list>, starships <list>
```

Applying Multiple Conditions with filter()

You can **combine multiple conditions** using:

- A comma (,) – Equivalent to using & (logical AND).
- | – Logical OR, keeping rows that meet **at least one** condition.

```
starwars |>
  filter(species %in% c("Human", "Droid"), height < 175) |>
  print(n = 3)
```

```
# A tibble: 14 x 14
  name   height  mass hair_color skin_color eye_color birth_year sex   gender homeworld species films
  <chr>   <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>    <chr>   <chr> <list>
1 Luke~     172     77 blond       fair        blue          19 male   masculin~ Tatooine Human   <chr>
2 C-3PO      167     75 <NA>       gold        yellow        112 none   masculin~ Tatooine Droid   <chr>
3 R2-D2      96      32 <NA>       white, bl~ red           33 none   masculin~ Naboo   Droid   <chr>
# i 11 more rows
# i 2 more variables: vehicles <list>, starships <list>
```

This keeps only: Characters who are either Human or Droid And have a height less than 175 cm. Alternatively, using explicit logical operators:

```
starwars |> filter(species == "Human" | species == "Droid", height < 175)
```

filter() - common mistakes

- = vs == for equality tests
 - ▶ Wrong: `filter(species = "Human")` (one equal sign leads to error)
 - ▶ Right: `filter(species == "Human")`
- Using OR incorrectly
 - ▶ Wrong: `filter(species == "Human" | "Droid")` (all rows will end up being selected)
 - ▶ Right: `filter(species %in% c("Human", "Droid"))` or `filter(species == "Human" | species == "Droid")`

Filtering with Regular Expressions

We will cover **regular expressions (regex)** in a future lecture. For now, note that `str_detect()` (from `stringr`) helps filter text-based patterns.

```
library(stringr)
starwars |>
  filter(str_detect(name, "^F"))
```

```
# A tibble: 2 x 14
  name   height  mass hair_color skin_color eye_color birth_year sex   gender homeworld species films
  <chr>   <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>    <chr>   <chr> <list>
1 Fini~     170     NA blond     fair       blue        91 male   masculin Coruscant Human   <chr>
2 Finn       NA      NA black     dark       dark        NA male   masculin <NA>      Human   <chr>
# i 2 more variables: vehicles <list>, starships <list>
```

- `str_detect(name, "F")` returns TRUE for names that start with "F" (^ means "start of the string").
- `filter()` keeps only rows where this condition holds.

Combining dplyr Functions with the Pipe (|>)

You can **chain multiple dplyr functions** using the pipe (|>), making your code more readable and efficient.

The following code keeps only:

- Characters with **no hair** (`hair_color == "none"`)
- OR characters with **black eyes** (`eye_color == "black"`)

Then, it selects only the specified columns.

```
starwars |>  
  filter(hair_color == "none" | eye_color == "black") |>  
  select(name, species, homeworld, hair_color, eye_color)
```

```
# A tibble: 39 x 5  
  name      species    homeworld   hair_color eye_color  
  <chr>     <chr>       <chr>        <chr>      <chr>  
1 Darth Vader Human     Tatooine     none       yellow  
2 Greedo     Rodian     Rodia       <NA>       black  
3 IG-88      Droid      <NA>        none       red  
4 Bossk      Trandoshan Trandosha   none       red  
5 Lobot      Human     Bespin       none       blue  
6 Ackbar     Mon Calamari Mon Cala    none       orange  
7 Nien Nunb Sullustan  Sullust     none       black  
8 Nute Gunray Neimodian Cato Neimoidia none       red  
9 Jar Jar Binks Gungan Naboo      none       orange
```

Sorting Rows with arrange()

The `arrange()` function sorts rows based on a variable. By default, it sorts in ascending order. Use `desc()` to sort in descending order.

The following first sorts characters by birth year (newest to oldest), then by mass (ascending) within each birth year.

```
starwars |>
  select(name, birth_year, height, mass) |>
  arrange(desc(birth_year), mass)
```

```
# A tibble: 87 x 4
  name          birth_year height   mass
  <chr>        <dbl>     <int>   <dbl>
1 Yoda           896       66     17
2 Jabba Desilijic Tiure  600      175    1358
3 Chewbacca      200      228    112
4 C-3PO          112      167     75
5 Dooku          102      193     80
6 Ki-Adi-Mundi  92       198     82
7 Qui-Gon Jinn   92       193     89
8 Finis Valorum  91       170     NA
9 Palpatine      82       170     75
10 Cliegg Lars    82      183     NA
# i 77 more rows
```

Copyright Miles Chen. For personal use only. Do not distribute.

Selecting Rows by Position with slice()

The slice() function selects rows based on **row number**.

The following line keeps only **rows 5 through 10**, maintaining the original order.

```
starwars |>  
  slice(5:10)
```

```
# A tibble: 6 x 14  
  name   height  mass hair_color skin_color eye_color birth_year sex   gender homeworld species films  
  <chr>   <int> <dbl> <chr>     <chr>      <chr>       <dbl> <chr> <chr>    <chr> <chr> <list>  
1 Leia~    150     49 brown    light     brown        19 fema~ femin~ Alderaan Human  <chr>  
2 Owen~    178    120 brown, gr~ light     blue         52 male  mascul~ Tatooine Human  <chr>  
3 Beru~    165     75 brown    light     blue        47 fema~ femin~ Tatooine Human  <chr>  
4 R5-D4    97      32 <NA>    white, red red        NA none  mascul~ Tatooine Droid   <chr>  
5 Bigg~    183     84 black    light     brown        24 male  mascul~ Tatooine Human  <chr>  
6 Obi-~    182     77 auburn, w~ fair     blue-gray    57 male  mascul~ Stewjon Human  <chr>  
# i 2 more variables: vehicles <list>, starships <list>
```

Randomly Sampling Rows with slice_sample()

The slice_sample() function selects a **random sample** of rows.

```
starwars |> slice_sample(n = 5)
```

```
# A tibble: 5 x 14
  name   height  mass hair_color skin_color eye_color birth_year sex   gender homeworld species films
  <chr>    <int> <dbl> <chr>       <chr>       <dbl> <chr> <chr> <chr> <chr> <chr> <list>
1 Shaa~     178     57 none      red, blue~ black           NA fema~ femin~ Shili   Togruta <chr>
2 Dormé     165     NA brown    light      brown           NA fema~ femin~ Naboo  Human   <chr>
3 Finn      NA      NA black    dark       dark            NA male  mascul~ <NA>   Human   <chr>
4 Mon ~     150     NA auburn   fair       blue            48 fema~ femin~ Chandrila Human <chr>
5 Obi~~     182     77 auburn, w~ fair      blue-gray        57 male  mascul~ Stewjon Human <chr>
# i 2 more variables: vehicles <list>, starships <list>
```

This is useful when:

- Previewing **random** portions of a dataset.
- Generating a random subset for **testing or analysis**.

Selecting with slice_min() and slice_max()

The `slice_min()` and `slice_max()` functions select rows with the **smallest** or **largest** values for a given column.

```
starwars |> slice_max(mass, n = 3)
```

```
# A tibble: 3 x 14
  name   height  mass hair_color skin_color eye_color birth_year sex   gender homeworld species films
  <chr>   <int> <dbl> <chr>       <chr>       <dbl> <chr> <chr>   <chr>   <chr>   <list>
1 Jabb~     175   1358 <NA>       green-tan~ orange        600 herm~ mascul~ Nal Hutta Hutt   <chr>
2 Grie~     216    159 none       brown, wh~ green, y~          NA male  mascul~ Kalee   Kaleesh <chr>
3 IG-88     200    140 none       metal      red           15 none  mascul~ <NA>     Droid   <chr>
# i 2 more variables: vehicles <list>, starships <list>
```

This is equivalent to:

```
starwars |> arrange(desc(mass)) |> head(3)
```

- `slice_max(mass, n = 3)` directly finds the top 3 characters with highest mass values.
- `slice_min()` works the same way but finds the smallest values.

distinct() - unique rows or combos

distinct() removes duplicate rows and returns only unique rows. If multiple columns are selected, then it returns only unique combos of specified columns.

```
# Remove duplicate rows (if any)
starwars |>
  distinct()
```

```
# A tibble: 87 x 14
  name      height  mass hair_color skin_color eye_color birth_year sex   gender homeworld species
  <chr>     <int> <dbl> <chr>       <chr>       <chr>       <dbl> <chr> <chr> <chr>       <chr>
1 Luke Skywalker 172    77 blond      fair        blue         19   male  masculin Tatooine Human
2 C-3PO          167    75 <NA>       gold        yellow       112  none  masculin Tatooine Droid 
3 R2-D2          96     32 <NA>       white, bl~ red         33   none  masculin Naboo   Droid 
4 Darth Vader    202    136 none       white        yellow       41.9 male  masculin Tatooine Human
5 Leia Organa    150    49 brown      light       brown        19   femal feminin Alderaan Human
6 Owen Lars      178    120 brown, gr~ light       blue         52   male  masculin Tatooine Human
7 Beru Whitesun 165    75 brown      light       blue         47   femal feminin Tatooine Human
8 R5-D4          97     32 <NA>       white, red red        NA   none  masculin Tatooine Droid 
9 Biggs Darkley  183    84 black      light       brown        24   male  masculin Tatooine Human
10 Obi-Wan Kenobi 182    77 auburn, w~ fair        blue-gray     57   male  masculin Stewjon Human
# i 77 more rows
# i 3 more variables: films <list>, vehicles <list>, starships <list>
```

```
starwars |>  
distinct(species)
```

```
# A tibble: 38 x 1  
  species  
  <chr>  
1 Human  
2 Droid  
3 Wookiee  
4 Rodian  
5 Hutt  
6 <NA>  
7 Yoda's species  
8 Trandoshan  
9 Mon Calamari  
10 Ewok  
# i 28 more rows
```

```
starwars |>  
distinct(sex, gender)
```

```
# A tibble: 6 x 2  
  sex           gender  
  <chr>         <chr>  
1 male          masculine  
2 none          masculine  
3 female        feminine  
4 hermaphroditic masculine  
5 <NA>          <NA>  
6 none          feminine
```

Keep other columns for the first occurrence:

```
starwars |>  
  distinct(sex, gender, .keep_all = TRUE)
```

```
# A tibble: 6 x 14  
  name    height  mass hair_color skin_color eye_color birth_year sex   gender homeworld species films  
  <chr>   <int> <dbl> <chr>     <chr>      <chr>       <dbl> <chr> <chr> <chr> <chr> <chr> <list>  
1 Luke~     172     77 blond     fair       blue        19 male   masculin~ Tatooine Human  <chr>  
2 C-3PO     167     75 <NA>     gold       yellow      112 none   masculin~ Tatooine Droid  <chr>  
3 Leia~     150     49 brown    light       brown       19 feminin~ feminin~ Alderaan Human  <chr>  
4 Jabb~     175    1358 <NA>     green-tan~ orange      600 herm~ masculin~ Nal Hutta Hutt  <chr>  
5 Jek ~     180     110 brown    fair       blue        NA <NA>  <NA>   Bestine ~ <NA>  <chr>  
6 R4-P~     96      NA none     silver, r~ red, blue     NA none   feminin~ <NA>   Droid  <chr>  
# i 2 more variables: vehicles <list>, starships <list>
```

Section 3

making new columns

mutate() - add or transform columns

mutate() computes new columns from existing ones. It keeps all existing columns by default, and puts the newly created column at the end.

```
flights |>  
  mutate(  
    gain = dep_delay - arr_delay,  
    speed_mph = distance / (air_time / 60)  
  )
```

```
# A tibble: 435,352 x 21  
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time arr_delay carrier  
  <int> <int> <int>     <int>       <int>      <dbl>     <int>       <int>      <dbl> <chr>  
1 2023     1     1        1            2038      203       328          3     205  UA  
2 2023     1     1       18            2300       78       228         135     53  DL  
3 2023     1     1       31            2344       47       500         426     34  B6  
4 2023     1     1       33            2140      173       238         2352    166  B6  
5 2023     1     1       36            2048      228       223         2252    211  UA  
6 2023     1     1      503            500        3       808         815     -7  AA  
7 2023     1     1      520            510       10       948         949     -1  B6  
8 2023     1     1      524            530       -6       645         710    -25  AA  
9 2023     1     1      537            520       17       926         818     68  UA  
10 2023    1     1      547            545        2       845         852     -7  NK  
# i 435,342 more rows
```

option `.before` will place the new columns before the specified columns. In general, it's fine to have the newly created column appear as the last column in the data, but for the lecture, it's helpful to see the new columns up front.

```
flights |>
  mutate(
    gain = dep_delay - arr_delay,
    speed_mph = distance / (air_time / 60),
    .before = everything())
```

A tibble: 435,352 x 21

	gain	speed_mph	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time	sched_arr_time	
	<dbl>	<dbl>	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>	<int>	<int>
1	-2	409.	2023	1	1	1	2038	203	328	3	
2	25	422.	2023	1	1	18	2300	78	228	135	
3	13	498.	2023	1	1	31	2344	47	500	426	
4	7	353.	2023	1	1	33	2140	173	238	2352	
5	17	366	2023	1	1	36	2048	228	223	2252	
6	10	423.	2023	1	1	503	500	3	808	815	
7	11	492.	2023	1	1	520	510	10	948	949	
8	19	363.	2023	1	1	524	530	-6	645	710	
9	-51	326.	2023	1	1	537	520	17	926	818	
10	9	407.	2023	1	1	547	545	2	845	852	

i 435,342 more rows

i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>, origin <chr>,
i dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dttm>

Use newly created columns within the same `mutate()` call:

```
flights |>
  mutate(
    gain = dep_delay - arr_delay,
    gain_per_hour = gain / (air_time / 60),
    .before = everything()
  )
```

```
# A tibble: 435,352 x 21
   gain gain_per_hour year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
   <dbl>        <dbl> <int> <int> <int>      <int>        <dbl>     <int>        <dbl>     <int>
1    -2       -0.327  2023     1     1       1          2038      203       328          3
2    25        13.9   2023     1     1      18          2300       78       228        135
3    13        4.11   2023     1     1      31          2344       47       500        426
4     7        3.89   2023     1     1      33          2140      173       238        2352
5    17        12.8   2023     1     1      36          2048      228       223        2252
6    10        3.90   2023     1     1      503         500        3       808        815
7    11        3.44   2023     1     1      520         510       10       948        949
8    19        9.58   2023     1     1      524         530       -6       645        710
9   -51       -11.9  2023     1     1      537         520       17       926        818
10     9        3.44   2023     1     1      547         545        2       845        852
# i 435,342 more rows
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>, origin <chr>,
# dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dttm>
```

Vectorized conditionals - use `if_else(condition, result_if_true, result_if_false)`:

```
flights |>
  mutate(
    long_delay = if_else(arr_delay > 60, "> 1 hr", " 1 hr"),
    .before = everything()
  )
```

A tibble: 435,352 x 20

```
long_delay  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time arr_delay
<chr>       <int> <int> <int>     <int>        <dbl>    <int>        <int>        <dbl>
1 > 1 hr      2023     1     1      1           2038      203      328          3      205
2  1 hr       2023     1     1     18           2300      78      228         135      53
3  1 hr       2023     1     1     31           2344      47      500         426      34
4 > 1 hr      2023     1     1     33           2140      173      238         2352     166
5 > 1 hr      2023     1     1     36           2048      228      223         2252     211
6  1 hr       2023     1     1     503          500       3      808         815      -7
7  1 hr       2023     1     1     520          510       10     948         949      -1
8  1 hr       2023     1     1     524          530      -6      645         710     -25
9 > 1 hr      2023     1     1     537          520       17     926         818      68
10 1 hr       2023    1     1     547          545       2      845         852      -7
```

i 435,342 more rows

i 10 more variables: carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dttm>

Multiple cases, use case_when():

```
flights |>
  mutate(
    delay_band = case_when(
      arr_delay <= 0 ~ "early/on time",
      arr_delay <= 30 ~ "0-30",
      arr_delay <= 120 ~ "31-120",
      TRUE ~ ">120"
    ) ,
    .before = everything()
  )
```

```
# A tibble: 435,352 x 20
  delay_band   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time arr_delay
  <chr>       <int> <int> <int>   <int>          <int>     <dbl>    <int>          <int>        <dbl>
1 >120         2023     1     1      1           2038      203     328            3        205
2 31-120       2023     1     1     18           2300      78     228           135        53
3 31-120       2023     1     1     31           2344      47     500           426        34
4 >120         2023     1     1     33           2140      173    238          2352       166
5 >120         2023     1     1     36           2048      228    223          2252       211
6 early/on t~  2023     1     1    503           500       3     808           815       -7
7 early/on t~  2023     1     1    520           510      10     948           949       -1
8 early/on t~  2023     1     1    524           530      -6     645           710      -25
9 31-120       2023     1     1    537           520      17     926           818       68
10 early/on t~ 2023     1     1    547           545       2     845           852       -7
```

New Columns will have the same number of rows

Important: When using `mutate()`, the new column will have an entry for every row in the data. This means that summary functions like `mean` will result in the same value duplicated across all rows.

Example: Adding a column for the mean of `mass` duplicates the value for all rows

```
flights |>
  filter(month == 1, day == 1) |>
  select(carrier, flight, dep_time:dep_delay) |>
  mutate(mean_dep_delay = mean(dep_delay, na.rm = TRUE))
```

```
# A tibble: 865 x 6
  carrier flight dep_time sched_dep_time dep_delay mean_dep_delay
  <chr>    <int>    <int>        <int>     <dbl>            <dbl>
1 UA        628       1          2038      203        17.8
2 DL        393      18          2300       78        17.8
3 B6        371      31          2344       47        17.8
4 B6       1053      33          2140      173        17.8
5 UA        219      36          2048      228        17.8
6 AA        499      503          500        3        17.8
7 B6        996      520          510       10        17.8
8 AA        981      524          530       -6        17.8
9 UA       1053      524          530       17        17.8
```

Useful Functions for `mutate()`

Here are some handy functions that work well inside `mutate()`:

- Row-wise (Element-wise) Functions
 - ▶ `pmin()` / `pmax()` – Find the minimum/maximum **for each row** across multiple columns.
- Cumulative Functions
 - ▶ `cummin()` / `cummax()` – Keep track of the **minimum/maximum seen so far**.
 - ▶ `cumsum()` / `cumprod()` – Compute the **cumulative sum/product**.
 - ▶ `cummean()` – Compute the **cumulative mean**.
- Logical Functions
 - ▶ `between(x, a, b)` – Checks if `x` is **between a and b**.
- Offset Functions
 - ▶ `lead(x, n)` / `lag(x, n)` – Shift values **forward** (`lead()`) or **backward** (`lag()`).
 - ▶ `ntile(x, n)` – Splits a column into **n equal-sized bins**.

mutate() Examples

```
flights |>
  filter(month == 1, day == 1) |>
  select(carrier, dep_delay, arr_delay) |>
  mutate(
    cummin_arr_del = cummin(arr_delay), # Keeps track of the smallest mass seen so far.
    pmin_delay = pmin(dep_delay, arr_delay), # Element-wise min between mass & birth_year
    lag2 = lag(pmin_delay, 2) # Shift massyear_pmin values back by 2 rows
  )
```

```
# A tibble: 865 x 6
  carrier dep_delay arr_delay cummin_arr_del pmin_delay lag2
  <chr>     <dbl>     <dbl>          <dbl>        <dbl>   <dbl>
1 UA         203       205          205        203      NA
2 DL          78        53           53        53      NA
3 B6          47        34           34        34      203
4 B6         173       166          34        166      53
5 UA         228       211          34        211      34
6 AA           3        -7           -7        -7      166
7 B6          10        -1           -7        -1      211
8 AA          -6       -25          -25       -25      -7
9 UA          17        68           -25       17      -1
10 NK          2        -7           -25       -7      -25
# i 855 more rows
```

Copyright Miles Chen. For personal use only. Do not distribute.

transmute() - mutate + keep only new columns

transmute() is like mutate(), but drops the original columns unless you explicitly keep them.

```
flights |>  
  transmute(  
    tailnum,  
    gain = dep_delay - arr_delay,  
    speed_mph = distance / (air_time / 60)  
  )
```

```
# A tibble: 435,352 x 3  
  tailnum   gain speed_mph  
  <chr>     <dbl>     <dbl>  
1 N25201     -2     409.  
2 N830DN      25     422.  
3 N807JB      13     498.  
4 N265JB       7     353.  
5 N17730      17     366  
6 N925AN      10     423.  
7 N2043J       11     492.  
8 N918AN      19     363.  
9 N13113     -51     326.  
10 N912NK      9     407.  
# i 435,342 more rows
```

Copyright Miles Chen. For personal use only. Do not distribute.

relocate() - move columns without dropping

relocate() moves existing columns relative to others. The .before option in the earlier examples work by making a call to relocate

relocate() is great when you want to keep all columns but improve readability.

```
# Move dep/arr columns to the front  
flights |>  
  relocate(dep_time, arr_time)
```

```
# A tibble: 435,352 x 19  
  dep_time arr_time year month day sched_dep_time dep_delay sched_arr_time arr_delay carrier  
    <int>     <int> <int> <int> <int>      <dbl>     <int>      <dbl> <chr>  
1       1      328  2023     1     1        2038      203         3   205 UA  
2      18      228  2023     1     1        2300       78        135      53 DL  
3      31      500  2023     1     1        2344       47        426      34 B6  
4      33      238  2023     1     1        2140      173        2352     166 B6  
5      36      223  2023     1     1        2048      228        2252     211 UA  
6     503      808  2023     1     1        500        3        815      -7 AA  
7     520      948  2023     1     1        510       10        949      -1 B6  
8     524      645  2023     1     1        530       -6        710      -25 AA  
9     537      926  2023     1     1        520       17        818      68 UA  
10    547      845  2023     1     1        545        2        852      -7 NK
```

```
# Move after a reference column
flights |>
  relocate(dep_delay, arr_delay, .after = arr_time)

# A tibble: 435,352 x 19
   year month   day dep_time sched_dep_time arr_time dep_delay arr_delay sched_arr_time carrier
   <int> <int> <int>    <int>        <int>    <int>    <dbl>    <dbl>        <int> <chr>
1  2023     1     1       1        2038      328     203     205          3  UA
2  2023     1     1      18       2300      228      78      53         135  DL
3  2023     1     1      31       2344      500      47      34         426  B6
4  2023     1     1      33       2140      238     173     166        2352  B6
5  2023     1     1      36       2048      223     228     211        2252  UA
6  2023     1     1     503       500      808      3      -7         815  AA
7  2023     1     1     520       510      948     10      -1         949  B6
8  2023     1     1     524       530      645     -6     -25        710  AA
9  2023     1     1     537       520      926     17      68         818  UA
10 2023     1     1     547       545      845      2      -7        852  NK
# i 435,342 more rows
# i 9 more variables: flight <int>, tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>,
#   distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dttm>
```

```
# Move to the end
flights |>
  relocate(contains("delay"), .after = last_col())

# A tibble: 435,352 x 19
  year month   day dep_time sched_dep_time arr_time sched_arr_time carrier flight tailnum origin
  <int> <int> <int>    <int>          <int>    <int>          <int> <chr>   <int> <chr>   <chr>
1 2023     1     1        1            2038      328           3  UA       628 N25201 EWR
2 2023     1     1       18            2300      228          135  DL      393 N830DN JFK
3 2023     1     1       31            2344      500          426  B6      371 N807JB JFK
4 2023     1     1       33            2140      238          2352 B6     1053 N265JB JFK
5 2023     1     1       36            2048      223          2252 UA      219 N17730 EWR
6 2023     1     1      503            500       808          815  AA      499 N925AN EWR
7 2023     1     1       520            510       948          949  B6      996 N2043J JFK
8 2023     1     1       524            530       645          710  AA      981 N918AN EWR
9 2023     1     1       537            520       926          818  UA      206 N13113 EWR
10 2023    1     1      547            545       845          852  NK      225 N912NK EWR
# i 435,342 more rows
# i 8 more variables: dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,
#   time_hour <dttm>, dep_delay <dbl>, arr_delay <dbl>
```

Putting it together

`lubridate::make_date()` combines year, month, day information into a date-type variable, which R understands as dates.

Create derived variables, trim columns, and reorder for reporting:

```
flights |>
  mutate(
    gain = dep_delay - arr_delay,
    speed_mph = distance / (air_time / 60)
  ) |>
  transmute(
    date = lubridate::make_date(year, month, day),
    dep_time, arr_time, origin, dest, carrier,
    gain, speed_mph
  ) |>
  relocate(date, carrier, origin, dest, .before = dep_time)
```

```
# A tibble: 435,352 x 8
  date      carrier origin dest  dep_time arr_time   gain speed_mph
  <date>    <chr>   <chr> <chr>   <int>    <int> <dbl>    <dbl>
1 2023-01-01 UA       EWR    SMF        1      328    -2     409.
2 2023-01-01 DL       JFK    ATL       18      228    25     422.
3 2023-01-01 B6       JFK    BQN       31      500    13     498.
4 2023-01-01 B6       JFK    CHS       33      238     7     353.
5 2023-01-01 UA       EWR    DTW       36      223    17     366
6 2023-01-01 AA       EWR    MIA      503      808    10     423.
7 2023-01-01 B6       JFK    BQN      520      948    11     492.
8 2023-01-01 AA       EWR    ORD      524      645    19     363.
9 2023-01-01 UA       EWR    IAH      537      926   -51     326.
10 2023-01-01 NK      EWR    FLL      547      845     9     407.
```