

Data Wrangling in the Tidyverse Part 2


Stats 102A

Miles Chen

Department of Statistics and Data Science

UCLA

Resource: R For Data Science



Portions of this lecture are derived from the book. The book is Free to read:
<https://r4ds.had.co.nz/>

Section 1

Grouping Data

Summarizing Data with summarize()

The `summarize()` function (also spelled `summarise()`) is used to compute summary statistics for a dataset. It reduces multiple values into a single summary value per group or for the entire dataset.

```
starwars |>
  select(height, mass) |>
  summarize(
    avg_height = mean(height, na.rm = TRUE),
    var_height = var(height, na.rm = TRUE),
    avg_mass = mean(mass, na.rm = TRUE),
    min_height = min(height, na.rm = TRUE),
    max_mass = max(mass, na.rm = TRUE),
    count = n()
  )
```

```
# A tibble: 1 x 6
  avg_height var_height avg_mass min_height max_mass count
  <dbl>       <dbl>      <dbl>     <int>      <dbl>   <int>
1      175.      1209.      97.3       66      1358     87
```

- `mean()`, `var()`, `min()`, `max()` compute **summary statistics**.
- `n()` counts the **number of non-missing rows**.
- `na.rm = TRUE` ensures that missing values are ignored.

Creating Groups with group_by()

The `group_by()` function categorizes data into groups, allowing you to perform computations within each group.

```
starwars |>  
  group_by(species) |>  
  select(name, height, mass, species)
```

```
# A tibble: 87 x 4  
# Groups:   species [38]  
  name           height   mass species  
  <chr>         <int>   <dbl> <chr>  
1 Luke Skywalker     172     77 Human  
2 C-3PO              167     75 Droid  
3 R2-D2              96      32 Droid  
4 Darth Vader        202    136 Human  
5 Leia Organa         150     49 Human  
6 Owen Lars           178    120 Human  
7 Beru Whitesun Lars  165     75 Human  
8 R5-D4              97      32 Droid  
9 Biggs Darklighter   183     84 Human  
10 Obi-Wan Kenobi    182     77 Human  
# i 77 more rows
```

This does **not** change the data directly but sets up groups for further operations.
Copyright Miles Chen. For personal use only. Do not distribute.

Combining group_by() and summarize()

The real power of `group_by()` comes when it is combined with `summarize()`.

The following computes **mean**, **standard deviation**, and **count** for each species.

```
starwars |>
  group_by(species) |>
  summarize(
    mean_ht = mean(height, na.rm = TRUE),
    sd_ht = sd(height, na.rm = TRUE),
    mean_mass = mean(mass, na.rm = TRUE),
    sd_mass = sd(mass, na.rm = TRUE),
    count = n()
  )
```

```
# A tibble: 38 x 6
  species   mean_ht   sd_ht   mean_mass   sd_mass   count
  <chr>     <dbl>     <dbl>     <dbl>     <dbl>     <int>
1 Aleena      79       NA        15       NA        1
2 Besalisk    198      NA       102       NA        1
3 Cerean      198      NA        82       NA        1
4 Chagrian    196      NA       NaN       NA        1
5 Clawdite    168      NA        55       NA        1
6 Droid       131.     49.1     69.8     51.0      6
7 Dug        112       NA        40       NA        1
8 Ewok        88        NA        20       NA        1
```

Filtering and Sorting Grouped Summaries

We can filter and arrange grouped summaries to focus on meaningful results.

```
starwars |>
  group_by(species) |>
  summarize(
    mean_ht = mean(height, na.rm = TRUE),
    sd_ht = sd(height, na.rm = TRUE),
    mean_mass = mean(mass, na.rm = TRUE),
    sd_mass = sd(mass, na.rm = TRUE),
    count = n()
  ) |>
  filter(count > 1) |> # Keep species with more than 1 character
  arrange(desc(count)) |> # Sort by count in descending order
  print(n = 6)
```

```
# A tibble: 9 x 6
  species  mean_ht  sd_ht  mean_mass  sd_mass  count
  <chr>     <dbl>   <dbl>     <dbl>    <dbl>   <int>
1 Human      178    12.0      81.3    19.3     35
2 Droid       131.   49.1      69.8    51.0      6
3 <NA>        175    12.4      81       31.2     4
4 Gungan      209.   14.2      74       11.3     3
5 Kaminoan    221    11.3      88       NA       2
6 Mirialian   168    2.88      68.1    4.38     2
```

group_by() + mutate()

mutate() can be combined with group_by() to compute new variables **within each group**.

```
starwars |>
  filter(species %in% c("Human", "Droid") | is.na(species)) |>
  select(name, species, height) |>
  group_by(species) |>
  mutate(z_height = (height - mean(height, na.rm = TRUE)) / sd(height, na.rm = TRUE)) |>
  print(n = 3)
```

```
# A tibble: 45 x 4
# Groups:   species [3]
  name      species height z_height
  <chr>     <chr>    <int>    <dbl>
1 Luke Skywalker Human     172    -0.498
2 C-3PO        Droid     167     0.728
3 R2-D2        Droid      96    -0.716
# i 42 more rows
```

- The **z-score** of height is computed **within each species**.
- A character's height is compared **to the mean height of their species**.

C-3PO is **above average** among droids but **below average** when compared to all characters.

Computing Without group_by()

If we **do not** use `group_by()`, `mutate()` computes statistics using **all characters** instead of grouping by `species`.

```
starwars |>
  filter(species %in% c("Human", "Droid") | is.na(species)) |>
  select(name, species, height) |>
  # group_by(species) |>
  mutate(z_height = (height - mean(height, na.rm = TRUE)) / sd(height, na.rm = TRUE)) |>
  print(n = 3)
```

```
# A tibble: 45 x 4
  name      species height z_height
  <chr>     <chr>    <int>   <dbl>
1 Luke Skywalker Human     172    0.0123
2 C-3PO        Droid     167   -0.188
3 R2-D2        Droid      96   -3.03
# i 42 more rows
```

- Without `group_by()`, **all characters** are used to compute `mean(height)`.
- C-3PO's z-score now compares his height **to all characters**, not just droids.

Grouping by Multiple Variables

You can provide `group_by()` with **multiple variables**, which creates a **hierarchical grouping**. This allows summarization at different levels.

```
toy_cases <- read_csv("https://raw.githubusercontent.com/rstudio/EDAWR/master/data-raw/toyb.csv")
print(toy_cases)
```

```
# A tibble: 12 x 4
  country     year   sex   cases
  <chr>      <dbl> <chr>  <dbl>
1 Afghanistan 1999 female    1
2 Afghanistan 1999 male     1
3 Afghanistan 2000 female    1
4 Afghanistan 2000 male     1
5 Brazil       1999 female    2
6 Brazil       1999 male     2
7 Brazil       2000 female    2
8 Brazil       2000 male     2
9 China        1999 female    3
10 China       1999 male     3
11 China       2000 female    3
12 China       2000 male     3
```

Multiple group_by() + summarize()

You can **group by more than one variable** to create hierarchical summaries.

```
summary1 <-
  toy_cases |>
  group_by(country, year) |>
  summarize(cases = sum(cases))
print(summary1)
```

```
# A tibble: 6 x 3
# Groups:   country [3]
  country     year cases
  <chr>       <dbl> <dbl>
1 Afghanistan 1999    2
2 Afghanistan 2000    2
3 Brazil      1999    4
4 Brazil      2000    4
5 China       1999    6
6 China       2000    6
```

Aggregating Across Grouped Summaries

We can further summarize across **higher levels of grouping**.

```
summary2 <- summary1 |> summarize(cases = sum(cases))
summary2
```

```
# A tibble: 3 x 2
  country     cases
  <chr>       <dbl>
1 Afghanistan     4
2 Brazil           8
3 China            12
```

```
summary3 <- summary2 |> summarize(cases = sum(cases))
summary3
```

```
# A tibble: 1 x 1
  cases
  <dbl>
1     24
```

Each summarize operation collapses another level of grouping, aggregating data progressively.

Changing the Order of group_by()

Changing the order of variables in group_by() affects how data is grouped and summarized.

Example: Grouping by year First, Then country

```
summary_a <- toy_cases |> group_by(year, country) |>
  summarize(cases = sum(cases))
print(summary_a)
```

```
# A tibble: 6 x 3
# Groups:   year [2]
  year country     cases
  <dbl> <chr>      <dbl>
1 1999 Afghanistan     2
2 1999 Brazil           4
3 1999 China             6
4 2000 Afghanistan     2
5 2000 Brazil           4
6 2000 China             6
```

Aggregating Across Reordered Groups

```
summary_b <- summary_a |> summarize(cases = sum(cases))
summary_b
```

```
# A tibble: 2 x 2
  year   cases
  <dbl>   <dbl>
1 1999     12
2 2000     12
```

```
summary_c <- summary_b |> summarize(cases = sum(cases))
summary_c
```

```
# A tibble: 1 x 1
  cases
  <dbl>
1     24
```

This progressively collapses the data one level at a time.

Section 2

Merging Tables

Two-table verbs

Two-table verbs with dplyr

When working with multiple datasets, you'll often need to combine them to get insights from multiple sources. The **dplyr** package provides several useful functions (verbs) to combine data:

- **Mutating joins:** Add new columns to one table based on matching rows from another.
- **Filtering joins:** Keep rows based on matches found (or not found) in another table.
- **Set operations:** Treat observations as elements of mathematical sets, allowing intersection, union, and differences.

Note: If you're familiar with SQL, these joins will look very similar.

For additional details, see the official documentation:

<https://dplyr.tidyverse.org/articles/two-table.html>

Vocabulary

- **Fact table:** many rows of events (e.g., flights)
- **Dimension table:** small lookup/labels (e.g., airlines, airports)
- **Primary key:** unique in dimension
- **Foreign key:** repeats in fact
- **Join order mantra:** *Start from the table you want rows from; join in what you need.*

Example Data Tables

We'll use two small "toy" tables to illustrate these concepts clearly:

```
people <- tibble(  
  name = c("Adam", "Betty", "Carl", "Doug"),  
  state = c("CA", "CA", "NY", "TX")  
)  
  
states <- tibble(  
  abbreviation = c("CA", "NY", "WA"),  
  state_name = c("California", "New York", "Washington")  
)
```

left_join()

people table		states table	
	name state		abbreviation state_name
1	Adam CA	1	CA California
2	Betty CA	2	NY New York
3	Carl NY	3	WA Washington
4	Doug TX		

`left_join()` takes all the values in the **left** table and adds columns from the right table by matching values using a column that exists in both tables. Values that do not exist in the other table have NA returned.

```
people |>  
  left_join(states, by = join_by(state == abbreviation))
```

```
# A tibble: 4 x 3  
  name   state state_name  
  <chr>  <chr> <chr>  
1 Adam    CA    California  
2 Betty   CA    California  
3 Carl    NY    New York  
4 Doug    TX    <NA>
```

Result: All names remain, but Doug (TX) has no matching `state_name`, so gets NA.

right_join()

people table		states table	
	name state		abbreviation state_name
1	Adam CA	1	CA California
2	Betty CA	2	NY New York
3	Carl NY	3	WA Washington
4	Doug TX		

`right_join()` is similar to `left_join` except it keeps all the rows in the `right` table, adding matching columns from the left table. Non-matching rows receive NA. In general, we recommend using only `left_join()` and switching the order of tables instead.

```
people |>  
  right_join(states, by = join_by(state == abbreviation))
```

```
# A tibble: 4 x 3  
  name state state_name  
  <chr> <chr> <chr>  
1 Adam  CA    California  
2 Betty CA    California  
3 Carl  NY    New York  
4 <NA>  WA    Washington
```

Result: All states remain, but “Washington” (WA) has no matching person, so gets NA.
Copyright Miles Chen. For personal use only. Do not distribute.

left_join() with table order switched

people table		states table	
	name state		abbreviation state_name
1	Adam CA	1	CA California
2	Betty CA	2	NY New York
3	Carl NY	3	WA Washington
4	Doug TX		

Note that the rows are the same, but the column orders are different. Also notice that the names in the by = argument are swapped.

```
states |>  
  left_join(people, by = join_by(abbreviation == state))
```

```
# A tibble: 4 x 3  
  abbreviation state_name name  
  <chr>        <chr>      <chr>  
1 CA           California Adam  
2 CA           California Betty  
3 NY           New York   Carl  
4 WA           Washington <NA>
```

inner_join()

people table		states table	
	name state		abbreviation state_name
1	Adam CA	1	CA California
2	Betty CA	2	NY New York
3	Carl NY	3	WA Washington
4	Doug TX		

inner_join() keeps only rows that have values that exist in **both** tables. You can think of this as the intersection.

```
people |>  
  inner_join(states, by = join_by(state == abbreviation))
```

```
# A tibble: 3 x 3  
  name   state state_name  
  <chr>  <chr> <chr>  
1 Adam   CA    California  
2 Betty  CA    California  
3 Carl   NY    New York
```

Result: Doug (TX) and Washington (WA) are omitted since no matching pair exists.

full_join()

people table		states table	
	name state		abbreviation state_name
1	Adam CA	1	CA California
2	Betty CA	2	NY New York
3	Carl NY	3	WA Washington
4	Doug TX		

full_join() keeps **all** rows from both tables. Rows without matches receive NA. You can think of this as the union.

```
people |>  
  full_join(states, by = join_by(state == abbreviation))
```

# A tibble: 5 x 3		
	name state	state_name
	<chr>	<chr>
1	Adam CA	California
2	Betty CA	California
3	Carl NY	New York
4	Doug TX	<NA>
5	<NA> WA	Washington

Result: All rows from both tables appear, with NA filling unmatched fields.

Controlling Matching Columns (by argument)

Depending on the tables, the join operation can match tables on different variables.

In the previous examples, we used a named character vector `by = join_by(state == abbreviation)` specifying the name in the left table that matches the name in the right table.

Options for joining tables

- `by = join_by(x)` when names match
- `by = join_by(x == y)` when names differ
- Although you can leave the `by =` argument blank, **you should always specify something for by =**

Non-unique matches (Cartesian explosions)

If the joining key is not unique, all possible combinations (Cartesian product) of matches will be included. Usually, this indicates an error in the data.

```
people <- tibble(  
  name = c("Adam", "Betty", "Carl", "Doug"),  
  state = c("CA", "CA", "NY", "TX")  
)  
places <- tibble(  
  abbreviation = c("CA", "NY", "WA", "CA"),  
  state_name = c("California", "New York", "Washington", "Canada")  
)
```

Non-unique matches

people table		places table	
	name state		abbreviation state_name
1	Adam CA	1	CA California
2	Betty CA	2	NY New York
3	Carl NY	3	WA Washington
4	Doug TX	4	CA Canada

```
people |> left_join(places, by = join_by(state == abbreviation))
```

```
# A tibble: 6 x 3
  name state state_name
  <chr> <chr> <chr>
1 Adam  CA   California
2 Adam  CA   Canada
3 Betty CA   California
4 Betty CA   Canada
5 Carl   NY   New York
6 Doug   TX   <NA>
```

- Adam and Betty from CA each appear twice, once for California, once for Canada.
- **Avoiding such issues:** Always ensure keys are unique unless duplicates are intentional.

De-duplicate first

You can use `distinct()` to get rid of duplicates. However, in this case, it means eliminating Canada (California comes before Canada in the alphabet).

```
places_dedup <- places |>  
  distinct(abbreviation, .keep_all = TRUE)  
  
people |> left_join(places_dedup, by = join_by(state == abbreviation))
```

```
# A tibble: 4 x 3  
  name state state_name  
  <chr> <chr> <chr>  
1 Adam   CA    California  
2 Betty  CA    California  
3 Carl   NY    New York  
4 Doug   TX    <NA>
```

Additional Join Operations

Filtering joins (`semi_join()` and `anti_join()`):

- `semi_join(x, y)`: Keep rows from x if they have matches in y. Does not add columns from y.
- `anti_join(x, y)`: Keep rows from x if they have no matches in y.

Example of semi_join():

```
people |>  
  semi_join(states, join_by(state == abbreviation))
```

```
# A tibble: 3 x 2
```

```
  name    state  
  <chr>   <chr>  
1 Adam    CA  
2 Betty   CA  
3 Carl    NY
```

- Returns Adam, Betty, Carl (matching states only, no added columns).

Example of anti_join():

```
people |>  
  anti_join(states, join_by(state == abbreviation))
```

```
# A tibble: 1 x 2
```

```
  name    state  
  <chr>   <chr>
```

```
1 Doug    TX
```

- Returns Doug (state TX with no match).

Set operations: Intersect / union / setdiff on rows

```
west_coast <- tibble(state = c("CA", "OR", "WA"))
people_states <- tibble(state = c("CA", "NY", "TX"))

intersect(people_states, west_coast)
```

```
# A tibble: 1 x 1
  state
  <chr>
1 CA
```

```
setdiff(people_states, west_coast)
```

```
# A tibble: 2 x 1
```

```
  state
```

```
  <chr>
```

```
1 NY
```

```
2 TX
```

```
union(people_states, west_coast)
```

```
# A tibble: 5 x 1
```

```
  state
```

```
  <chr>
```

```
1 CA
```

```
2 NY
```

```
3 TX
```

```
4 OR
```

```
5 WA
```

Section 3

Reshaping Tables

Section 4

Reshaping Data

Toy data sets

```
storms_url <- "https://raw.githubusercontent.com/rstudio/EDAWR/master/data-raw/storms.csv"
storms <- read_csv(storms_url, show_col_types = FALSE)
cases_url <- "https://raw.githubusercontent.com/rstudio/EDAWR/master/data-raw/cases.csv"
cases <- read_csv(cases_url, show_col_types = FALSE)
pollution_url <- "https://raw.githubusercontent.com/rstudio/EDAWR/master/data-raw/pollution.csv"
pollution <- read_csv(pollution_url, show_col_types = FALSE)
```

Storms Table

storms

```
# A tibble: 6 x 4
  storm    wind pressure date
  <chr>   <dbl>    <dbl> <date>
1 Alberto     110     1007 2000-08-03
2 Alex        45      1009 1998-07-27
3 Allison     65      1005 1995-06-03
4 Ana         40      1013 1997-06-30
5 Arlene       50      1010 1999-06-11
6 Arthur       45      1010 1996-06-17
```

- Each row represents a different tropical storm. The observational unit is a tropical storm.
- The variables are:
 - ▶ storm name
 - ▶ wind speed
 - ▶ air pressure
 - ▶ date
- We have one column for each variable
- This data is tidy as is. No reshaping needed.

mutate() works well

```
storms |> mutate(ratio = pressure / wind)
```


```
# A tibble: 6 x 5
```

	storm	wind	pressure	date	ratio
	<chr>	<dbl>	<dbl>	<date>	<dbl>
1	Alberto	110	1007	2000-08-03	9.15
2	Alex	45	1009	1998-07-27	22.4
3	Allison	65	1005	1995-06-03	15.5
4	Ana	40	1013	1997-06-30	25.3
5	Arlene	50	1010	1999-06-11	20.2
6	Arthur	45	1010	1996-06-17	22.4

ggplot() works well

With each variable in its own column, we can map the variables to different aesthetics.

```
library(ggrepel)
storms |> ggplot(aes(x = wind, y = pressure, color = date)) +
  geom_point(size = 5) +
  geom_text_repel(aes(label = storm))
```



Cases table

This is a table with fictional data regarding the cases of a certain infection in different years.

`cases`

```
# A tibble: 3 x 4
  country `2011` `2012` `2013`
  <chr>    <dbl>   <dbl>   <dbl>
1 FR        7000    6900    7000
2 DE        5800    6000    6200
3 US       15000   14000   13000
```

- Each row is a country. The observational unit is country.
- The variables are:
 - ▶ number of cases in 2011
 - ▶ number of cases in 2012
 - ▶ number of cases in 2013


mutate() works well

```
cases |> mutate(change_11_12 = `2012` - `2011`, change_12_13 = `2013` - `2012`, )  
  
# A tibble: 3 x 6  
  country `2011` `2012` `2013` change_11_12 change_12_13  
  <chr>    <dbl>   <dbl>   <dbl>      <dbl>       <dbl>  
1 FR        7000    6900    7000      -100        100  
2 DE        5800    6000    6200       200        200  
3 US       15000   14000   13000     -1000     -1000
```

However, plotting is limited

Because each column represents a different year, it is difficult to produce a plot that allows us to compare the different years against each other. `ggplot()` requires that only one column gets mapped to `y`. To see both 2011 and 2012 numbers, I would have to make separate plots.


```
cases |> ggplot(aes(x = country, y = `2011`)) +  
  geom_col() +  
  ggtitle("Cases for 2011")
```



Plotting is limited

Attempting to plot the different years by adding multiple `geom_col()` layers is possible, but results in a plot that is not very readable without additional work.

```
cases |> ggplot(aes(x = country)) +  
  geom_col(aes(y = `2011`), col = "red") +  
  geom_col(aes(y = `2012`), col = "blue") +  
  geom_col(aes(y = `2013`), col = "purple")
```



Pivoting to the rescue!

What I need to do is reshape or pivot the data so that each row represents the count of cases in a country for a given year.

I want the variables to be:

- country
- year
- count of cases

If the data is structured this way, I can map the year column to an aesthetic attribute.

pivot_longer()

To achieve the desired result, we use the function `pivot_longer()` because we want the resulting data set to be longer than the original data. (Older versions called this `gather()`). After pivoting the data, notice that we now have 9 rows - one for each country and year.

```
pivot_longer(cases,
             cols = "2011":"2013",
             names_to = "year",
             values_to = "cases")
```

```
# A tibble: 9 x 3
  country year  cases
  <chr>   <chr> <dbl>
1 FR      2011   7000
2 FR      2012   6900
3 FR      2013   7000
4 DE      2011   5800
5 DE      2012   6000
6 DE      2013   6200
7 US      2011  15000
8 US      2012  14000
9 US      2015  13000
```

The pivot_longer() function

```
pivot_longer(data = cases,  
            cols = "2011":"2013",  
            names_to = "year",  
            values_to = "cases")
```

The pivot_longer() function takes in a few arguments:

- data is the name of the data.frame or tibble that we will pivot
- cols are the names of the columns that will be pivoted. In this case, we want the columns named “2011” through “2013”. You can specify a range of column names with the : operator. Otherwise, you can provide a vector of column names
- names_to is a character string with what you want to call the resulting column of names. The former column names will be put into this column.
- values_to is a character string with what you want to call the resulting column of values. The former cell values will be put into this column.


The names in pivot_longer() are arbitrary

```
pivot_longer(cases,
             cols = "2011":"2013",
             names_to = "when it happened",
             values_to = "how many")
```

```
# A tibble: 9 x 3
  country `when it happened` `how many`
  <chr>    <chr>           <dbl>
1 FR       2011            7000
2 FR       2012            6900
3 FR       2013            7000
4 DE       2011            5800
5 DE       2012            6000
6 DE       2013            6200
7 US       2011           15000
8 US       2012           14000
9 US       2013           13000
```


Plotting is better with the data in this form

```
cases |>
  pivot_longer(cols = "2011":"2013",
               names_to = "year",
               values_to = "cases") |>
  ggplot(aes(x = country, y = cases, fill = year)) +
  geom_col(position = "dodge")
```



Another plot example

```
cases |>
  pivot_longer(cols = "2011":"2013",
               names_to = "year",
               values_to = "cases") |>
  ggplot(aes(x = year, y = cases, color = country)) +
  geom_point() + geom_path(aes(group = country))
```



Additional methods for selecting columns to pivot

```
cases |>
  pivot_longer(
    cols = starts_with("20"), # any column that starts with "20"
    names_to = "year",
    values_to = "cases"
)

cases |>
  pivot_longer(
    cols = -country, # everything except the country column
    names_to = "year",
    values_to = "cases"
)
```

The pollution table

Here is another table with fictional data regarding pollution measurements for different cities.

`pollution`


```
# A tibble: 6 x 3
  city      size   amount
  <chr>    <chr>   <dbl>
1 New York large     23
2 New York small    14
3 London    large    22
4 London    small    16
5 Beijing   large   121
6 Beijing   small    56
```

- Each row is a city-particle size combination. The observational unit is a pollution measurement for a specific city and particle size
- The variables are:
 - ▶ city
 - ▶ size
 - ▶ amount

Plotting is good with the current structure

I can map each column to an aesthetic attribute.

```
pollution |>  
  ggplot(aes(x = city, y = amount, fill = size)) +  
  geom_col(position = "dodge")
```



However, we can't do calculations with a single column

```
pollution
```

```
# A tibble: 6 x 3
  city      size   amount
  <chr>    <chr>   <dbl>
1 New York large     23
2 New York small    14
3 London   large     22
4 London   small     16
5 Beijing  large    121
6 Beijing  small     56
```

With the current structure, it's not possible (without a lot trouble) to calculate the difference between large and small particles in each city.

pivot_wider()

For this data, we need to use the function `pivot_wider()`. This function takes values in one column and pivots them across multiple columns (making the data wider). (Older versions called this `spread()`)

```
pollution |> pivot_wider(id_cols = city,
                           names_from = size,
                           values_from = amount)
```

```
# A tibble: 3 x 3
  city     large small
  <chr>    <dbl> <dbl>
1 New York     23    14
2 London       22    16
3 Beijing      121   56
```

pivot_wider()

```
pollution |> pivot_wider(id_cols = city,  
                           names_from = "size",  
                           values_from = "amount")
```

The `pivot_wider()` function takes in a few arguments:

- `data` is the name of the `data.frame` or `tibble` that we will pivot
- `id_cols` (optional) is the name of the column that identifies each observational unit
- `names_from` is the name of the column that has the names that will become column headers
- `values_from` is the name of the column that has the values

Calculations with `mutate()` for each city are now possible

```
pollution |>  
  pivot_wider(id_cols = city,  
              names_from = "size",  
              values_from = "amount") |>  
  mutate(difference = large - small,  
         ratio = large / small)
```

```
# A tibble: 3 x 5  
  city     large small difference ratio  
  <chr>    <dbl> <dbl>      <dbl> <dbl>  
1 New York     23     14          9   1.64  
2 London       22     16          6   1.38  
3 Beijing      121    56         65   2.16
```

Section 5

Warnings and Things to avoid

`pivot_longer()` duplicates any column that isn't pivoted

```
# A tibble: 3 x 4
  country `2011` `2012` `2013`
  <chr>    <dbl>   <dbl>   <dbl>
1 FR        7000    6900    7000
2 DE        5800    6000    6200
3 US       15000   14000   13000
```

Watch what happens if the columns I pivot are only “2012” and “2013”.

```
pivot_longer(cases,
             cols = "2012":"2013",
             names_to = "year",
             values_to = "cases")
```

pivot_longer() duplicates any column that isn't pivoted

The columns that are not pivoted are duplicated for the new rows created.

```
pivot_longer(cases,
             cols = "2012":"2013",
             names_to = "year",
             values_to = "cases")
```

```
# A tibble: 6 x 4
  country `2011` year  cases
  <chr>    <dbl> <chr> <dbl>
1 FR        7000  2012   6900
2 FR        7000  2013   7000
3 DE        5800  2012   6000
4 DE        5800  2013   6200
5 US       15000  2012  14000
6 US       15000  2013  13000
```

pivot_wider() is sensitive to spelling differences

If we have multiple spellings for New York (say New York and NYC), then each unique value gets its own row.

```
pollution2 <- pollution
pollution2[1,1] <- "NYC"
pollution2
```

```
# A tibble: 6 x 3
  city      size  amount
  <chr>    <chr>  <dbl>
1 NYC       large   23
2 New York small   14
3 London    large   22
4 London    small   16
5 Beijing   large  121
6 Beijing   small   56
```

Result

```
pivot_wider(pollution2,
            id_cols = city,
            names_from = "size",
            values_from = "amount")
```

```
# A tibble: 4 x 3
  city      large small
  <chr>    <dbl> <dbl>
1 NYC        23     NA
2 New York   NA     14
3 London     22     16
4 Beijing    121    56
```

Fill values

Sometimes you truly do have a scenario where you want to pivot wider and some entries do not exist. If you don't want NAs to show, you can specify a fill value.

Note: only use 0 if "missing" truly means zero. Otherwise leave NA or annotate the imputation.

```
pivot_wider(pollution2,
            id_cols = city,
            names_from = "size",
            values_from = "amount",
            values_fill = 0)
```

```
# A tibble: 4 x 3
  city      large small
  <chr>    <dbl> <dbl>
1 NYC        23     0
2 New York   0     14
3 London     22     16
4 Beijing    121    56
```

Another example of spelling differences

What will happen?

```
pollution2 <- pollution
pollution2[1,2] <- "LARGE"
pollution2
```

```
# A tibble: 6 x 3
  city      size   amount
  <chr>    <chr>   <dbl>
1 New York  LARGE     23
2 New York  small     14
3 London    large     22
4 London    small     16
5 Beijing   large    121
6 Beijing   small     56
```

Result

```
pivot_wider(pollution2,  
            id_cols = city,  
            names_from = "size",  
            values_from = "amount")
```

```
# A tibble: 3 x 4  
#>   city     LARGE small large  
#>   <chr>    <dbl> <dbl> <dbl>  
#> 1 New York    23    14    NA  
#> 2 London      NA    16    22  
#> 3 Beijing     NA    56   121
```

pivot_longer() and pivot_wider() are inverse operations

```
pollution
```

```
# A tibble: 6 x 3
  city    size   amount
  <chr>   <chr>   <dbl>
1 New York large     23
2 New York small    14
3 London   large    22
4 London   small    16
5 Beijing  large   121
6 Beijing  small    56

w <- pivot_wider(pollution, id_cols = city, names_from = "size", values_from = "amount")
w

# A tibble: 3 x 3
  city      large  small
  <chr>    <dbl> <dbl>
1 New York     23    14
2 London       22    16
3 Beijing      121   56
```

pivot_longer() and pivot_wider() are inverse operations

```
w

# A tibble: 3 x 3
  city      large  small
  <chr>     <dbl> <dbl>
1 New York    23    14
2 London      22    16
3 Beijing     121   56
pivot_longer(w, cols = "large":"small", names_to = "size", values_to = "amount")

# A tibble: 6 x 3
  city      size  amount
  <chr>     <chr> <dbl>
1 New York  large    23
2 New York  small    14
3 London    large    22
4 London    small    16
5 Beijing   large   121
6 Beijing   small    56
```

pivot_longer() and pivot_wider() are inverse operations

cases

```
# A tibble: 3 x 4
  country `2011` `2012` `2013`
  <chr>    <dbl>   <dbl>   <dbl>
1 FR        7000    6900    7000
2 DE        5800    6000    6200
3 US       15000   14000   13000
1 <- pivot_longer(cases, cols = "2011":"2013", names_to = "year", values_to = "count")
1
```

```
# A tibble: 9 x 3
  country year  count
  <chr>   <chr> <dbl>
1 FR      2011   7000
2 FR      2012   6900
3 FR      2013   7000
4 DE      2011   5800
5 DE      2012   6000
6 DE      2013   6200
7 US      2011  15000
8 US      2012  14000
9 US      2013  13000
```

pivot_longer() and pivot_wider() are inverse operations

```
1
```

```
# A tibble: 9 x 3
  country year  count
  <chr>   <chr> <dbl>
1 FR      2011   7000
2 FR      2012   6900
3 FR      2013   7000
4 DE      2011   5800
5 DE      2012   6000
6 DE      2013   6200
7 US      2011  15000
8 US      2012  14000
9 US      2013  13000
pivot_wider(l, names_from = "year", values_from = "count")
```

```
# A tibble: 3 x 4
  country `2011` `2012` `2013`
  <chr>    <dbl>  <dbl>  <dbl>
1 FR       7000   6900   7000
2 DE       5800   6000   6200
3 US      15000  14000  13000
```